



Java SE Beyond Basics: Generics, Annotation, Concurrency, and JMX

Sang Shin


sang.shin@sun.com

www.javapassion.com

Sun Microsystems Inc.

**UNLOCK
OPPORTUNITY**

What will you open?



SUN TECH DAYS 2006-2007
A Worldwide Developer Conference

Agenda

- Generics
- Annotation
- Concurrency
- JMX (Management & Monitoring)
- Performance
- Summary

Hands-on Labs

- Generics
 - > www.javapassion.com/handsonlabs/1110_tigergenerics.zip
- Annotation
 - > www.javapassion.com/handsonlabs/1111_javase5generics.zip
- Concurrency
 - > www.javapassion.com/handsonlabs/1108_javase5concurrency.zip
- JMX (Management & Monitoring)
 - > www.netbeans.org/kb/articles/jmx-tutorial.html

Generics

Sub-topics of Generics

- What is and why use Generics?
- Usage of Generics
- Generics and sub-typing
- Wildcard
- Type erasure
- Interoperability
- Creating your own Generic class

Generics:

What is it?

How do define it?

How to use it?

Why use it?

What is Generics?

- Generics provides abstraction over Types
 - > Classes, Interfaces and Methods can be **Parameterized** by **Types** (in the same way a Java type is parameterized by an instance of it)
- Generics makes **type safe code** possible
 - > If it compiles without any errors or warnings, then it **must not raise** any unexpected **ClassCastException** during runtime
- Generics provides increased readability
 - > Once you get used to it

Definition of a Generic Class: `LinkedList<E>`

- Definitions: `LinkedList<E>` has a type parameter `E` that represents the type of the elements stored in the linked list

```
public class LinkedList<E>
    extends AbstractSequentialList<E>
    implements List<E>, Queue<E>, Cloneable, java.io.Serializable{
    private transient Entry<E> header = new Entry<E>(null, null, null);
    private transient int size = 0;

    public E getFirst() {
        if (size==0) throw new NoSuchElementException();
        return header.next.element;
    }
```


Usage of Generic Class: LinkedList<Integer>

- Usage: Replace **type parameter** <E> with concrete **type argument**, like <Integer> or <String> or <MyType>
 - > LinkedList<Integer> can store only Integer or sub-type of Integer as elements


```
LinkedList<Integer> li =  
    new LinkedList<Integer>();  
li.add(new Integer(0));  
Integer i = li.iterator().next();
```

Example: Definition and Usage of Parameterized List interface

```
// Definition of the Generic'ized
// List interface
```

```
//
interface List<E>{
    void add(E x);
    Iterator<E> iterator();
    ...
}
```

Type parameter



```
// Usage of List interface with
// concrete type parameter, String
//
```

```
List<String> ls = new ArrayList<String>(10);
```

Type argument



Why Generics? Non-genericized Code is not Type Safe

```
// Suppose you want to maintain String
// entries in a Vector. By mistake,
// you add an Integer element. Compiler
// does not detect this. This is not
// type safe code.
```

```
Vector v = new Vector();
v.add(new String("valid string")); // intended
v.add(new Integer(4)); // unintended
```

```
// ClassCastException occurs during runtime
String s = (String)v.get(0);
```

Why Generics?

- Problem: Collection element types
 - > Compiler is unable to verify types of the elements
 - > Assignment must have type casting
 - > ClassCastException can occur during runtime
- Solution: Generics
 - > Tell the compiler the type of the collection
 - > Let the compiler do the casting
 - > Example: Compiler will check if you are adding Integer type entry to a String type collection
 - > Compile time detection of type mismatch

Generics:

Usage of Generics

Using Generic Classes: Example 1

- Instantiate a generic class to create type specific object
- In J2SE 5.0, all collection classes are rewritten to be generic classes

```
// Create a Vector of String type
Vector<String> vs = new Vector<String>();
vs.add(new Integer(5)); // Compile error!
vs.add(new String("hello"));
String s = vs.get(0);    // No casting needed
```

Using Generic Classes: Example 2

- Generic class can have multiple type parameters
- Type argument can be a custom type

```
// Create HashMap with two type parameters
HashMap<String, Mammal> map =
    new HashMap<String, Mammal>();
map.put("wombat", new Mammal("wombat"));

Mammal w = map.get("wombat");
```

Generics: Sub-typing

Generics and Sub-typing

- You can do this (using pre-J2SE 5.0 Java)
 - > `Object o = new Integer(5);`
- You can even do this (using pre-J2SE 5.0 Java)
 - > `Object[] or = new Integer[5];`
- So you would expect to be able to do this (Well, you can't do this!!!)
 - > `ArrayList<Object> ao = new ArrayList<Integer>();`
 - > This is counter-intuitive at the first glance

Generics and Sub-typing

- Why this compile error? It is because if it is allowed, ClassCastException can occur during runtime – **this is not type-safe**
 - > `ArrayList<Integer> ai = new ArrayList<Integer>();`
 - > `ArrayList<Object> ao = ai; // If it is allowed at compile time,`
 - > `ao.add(new Object());`
 - > `Integer i = ai.get(0); // This would result in`
`// runtime ClassCastException`
- So there is **no inheritance relationship between type arguments** of a generic class

Generics and Sub-typing

- The following code work
 - > `ArrayList<Integer> ai = new ArrayList<Integer>();`
 - > `List<Integer> li2 = new ArrayList<Integer>();`
 - > `Collection<Integer> ci = new ArrayList<Integer>();`
 - > `Collection<String> cs = new Vector<String>(4);`
- Inheritance relationship between generic classes themselves still exists

Generics and Sub-typing

- The following code work
 - > `ArrayList<Number> an = new ArrayList<Number>();`
 - > `an.add(new Integer(5));` // OK
 - > `an.add(new Long(1000L));` // OK
 - > `an.add(new String("hello"));` // compile error
- Entries in a collection maintain inheritance relationship

Generics: Wild card

Why Wildcards? Problem

- Consider the problem of writing a routine that prints out all the elements in a collection
- Here's how you might write it in an older version of the language (i.e., a pre-5.0 release):

```
static void printCollection(Collection c) {  
    Iterator i = c.iterator();  
    for (k = 0; k < c.size(); k++) {  
        System.out.println(i.next());  
    }  
}
```

Why Wildcards? Problem

- And here is a naive attempt at writing it using generics (and the new for loop syntax): **Well.. You can't do this!**

```
static void printCollection(Collection<Object> c) {  
    for (Object o : c)  
        System.out.println(o);  
}
```

```
public static void main(String[] args) {  
    Collection<String> cs = new Vector<String>();  
    printCollection(cs); // Compile error  
    List<Integer> li = new ArrayList<Integer>(10);  
    printCollection(li); // Compile error  
}
```

Why Wildcards? Solution

- Use Wildcard type argument `<?>`
- `Collection<?>` means **Collection of unknown type**
- Accessing entries of Collection of unknown type with **Object** type is safe

```
static void printCollection(Collection<?> c) {  
    for (Object o : c)  
        System.out.println(o);  
}  
  
public static void main(String[] args) {  
    Collection<String> cs = new Vector<String>();  
    printCollection(cs); // No Compile error  
    List<Integer> li = new ArrayList<Integer>(10);  
    printCollection(li); // No Compile error  
}
```


More on Wildcards

- You cannot access entries of Collection of unknown type other than **Object** type

```
static void printCollection(Collection<?> c) {  
    for (String o : c) // Compile error  
        System.out.println(o);  
}  
  
public static void main(String[] args) {  
    Collection<String> cs = new Vector<String>();  
    printCollection(cs); // No Compile error  
    List<Integer> li = new ArrayList<Integer>(10);  
    printCollection(li); // No Compile error  
}
```

More on Wildcards

- It isn't safe to add arbitrary objects to it however, since we don't know what the element type of `c` stands for, we cannot add objects to it.

```
static void printCollection(Collection<?> c) {  
    c.add(new Object()); // Compile time error  
    c.add(new String()); // Compile time error  
}  
  
public static void main(String[] args) {  
    Collection<String> cs = new Vector<String>();  
    printCollection(cs); // No Compile error  
    List<Integer> li = new ArrayList<Integer>(10);  
    printCollection(li); // No Compile error  
}
```

Bounded Wildcard

- If you want to bound the unknown type to be a subtype of another type, use Bounded Wildcard

```
static void printCollection(  
    Collection<? extends Number> c) {  
    for (Object o : c)  
        System.out.println(o);  
}  
  
public static void main(String[] args) {  
    Collection<String> cs = new Vector<String>();  
    printCollection(cs); // Compile error  
    List<Integer> li = new ArrayList<Integer>(10);  
    printCollection(li); // No Compile error  
}
```

Generics: Raw Type & Type Erasure

Raw Type

- Generic type instantiated with no type arguments
- Pre-J2SE 5.0 classes continue to function over J2SE 5.0 JVM as raw type

```
// Generic type instantiated with type argument  
List<String> ls = new LinkedList<String>();
```

```
// Generic type instantiated with no type  
// argument - This is Raw type  
List lraw = new LinkedList();
```

Type Erasure

- All generic type information is removed in the resulting byte-code after compilation
- So generic type information does not exist during runtime
- After compilation, they all share same class
 - > The class that represents `ArrayList<String>`, `ArrayList<Integer>` is the same class that represents `ArrayList`

Type Erasure Example Code: True or False?

```
ArrayList<Integer> ai = new ArrayList<Integer>();  
ArrayList<String> as = new ArrayList<String>();  
Boolean b1 = (ai.getClass() == as.getClass());  
System.out.println("Do ArrayList<Integer> and ArrayList<String> share  
    same class? " + b1);
```

Type-safe Code Again

- The compiler guarantees that either:
 - > the code it generates will be type-correct at run time, or
 - > it will output a warning (using Raw type) at compile time – in this case, you are responsible to make sure the warning is a benign one
- What is “type-safe code” again?
 - > If your code compiles without any compile errors and without warnings (or with warnings on safe operations), then you will never get a `ClassCastException` during runtime

Generics: Interoperability

What Happens to the following Code?

```
import java.util.LinkedList;
import java.util.List;

public class GenericsInteroperability {

    public static void main(String[] args) {

        List<String> ls = new LinkedList<String>();
        List lraw = ls;
        lraw.add(new Integer(4));
        String s = ls.iterator().next();
    }
}
```

Compilation and Running

- Compilation results in a warning message
 - > GenericsInteroperability.java uses unchecked or unsafe operations.
- Running the code
 - > ClassCastException

Generics: Creating Your Own Generic Class

Defining Your Own Generic Class

```
public class Pair<F, S> {  
    F first; S second;  
  
    public Pair(F f, S s) {  
        first = f; second = s;  
    }  
  
    public void setFirst(F f){  
        first = f;  
    }  
  
    public F getFirst(){  
        return first;  
    }  
  
    public void setSecond(S s){  
        second = s;  
    }  
  
    public S getSecond(){  
        return second;  
    }  
}
```

Using Your Own Generic Class

```
public class MyOwnGenericClass {  
    public static void main(String[] args) {  
        // Create an instance of Pair <F, S> class. Let's call it p1.  
        Number n1 = new Integer(5);  
        String s1 = new String("Sun");  
        Pair<Number,String> p1 = new Pair<Number,String>(n1, s1);  
        System.out.println("first of p1 (right after creation) = " + p1.getFirst());  
        System.out.println("second of p2 (right after creation) = " + p1.getSecond());  
  
        // Set internal variables of p1.  
        p1.setFirst(new Long(6L));  
        p1.setSecond(new String("rises"));  
        System.out.println("first of p1(after setting values) = " + p1.getFirst());  
        System.out.println("second of p1 (after setting values) = " + p1.getSecond());  
    }  
}
```

Annotation

Sub-topics of Annotations

- What is and Why annotation?
- How to define and use Annotations?
- 3 different kinds of Annotations
- Meta-Annotations

How Annotation Are Used?

- Annotations are used to affect the way programs are treated by tools and libraries
- Annotations are used by tools to produce derived files
 - > Tools: Compiler, IDE, Runtime tools
 - > Derived files : New Java code, deployment descriptor, class files

Ad-hoc Annotation-like Examples in pre-J2SE 5.0 Platform

- Ad-hoc Annotation-like examples in pre-J2SE 5.0 platform
 - > **Transient**
 - > **Serializable** interface
 - > **@deprecated**
 - > javadoc comments
 - > Xdoclet
- J2SE 5.0 Annotation provides a standard, general purpose, more powerful annotation scheme

Why Annotation?

- Enables “declarative programming” style
 - > Less coding since tool will generate the boiler plate code from annotations in the source code
 - > Easier to change
- Eliminates the need for maintaining "side files" that must be kept up to date with changes in source files
 - > Information is kept in the source file
 - > example) Eliminate the need of deployment descriptor

Annotation:
**How do you define &
use annotations?**

How to “Define” Annotation Type?

- Annotation type definitions are similar to normal Java **interface** definitions
 - > An at-sign (@) precedes the **interface** keyword
 - > **Each method declaration defines an element of the annotation type**
 - > Method declarations must not have any parameters or a throws clause
 - > Return types are restricted to primitives, String, Class, enums, annotations, and arrays of the preceding types
 - > Methods can have default values

Example: Annotation Type Definition

```
/**
```

```
 * Describes the Request-For-Enhancement(RFE) that led
```

```
 * to the presence of the annotated API element.
```

```
 */
```

```
public @interface RequestForEnhancement {
```

```
    int id();
```

```
    String synopsis();
```

```
    String engineer() default "[unassigned]";
```

```
    String date() default "[unimplemented]";
```

```
}
```

How To “Use” Annotation

- Once an annotation type is defined, you can use it to annotate declarations
 - > class, method, field declarations
- An annotation is a special kind of modifier, and can be used anywhere that other modifiers (such as public, static, or final) can be used
 - > By convention, annotations precede other modifiers
 - > Annotations consist of an at-sign (@) followed by an annotation type and a parenthesized list of element-value pairs

Example: Usage of Annotation

```
@RequestForEnhancement(  
    id      = 2868724,  
    synopsis = "Enable time-travel",  
    engineer = "Mr. Peabody",  
    date     = "4/1/3007"  
)  
public static void travelThroughTime(Date destination) {  
    ... }  
}
```

It is annotating **travelThroughTime** method

Annotation:

3 Types of Annotations (in terms of Sophistication)

3 Different Kinds of Annotations

- Marker annotation
- Single value annotation
- Normal annotation

Marker Annotation

- An annotation type with no elements
 - > Simplest annotation
- Definition

```
/**  
 * Indicates that the specification of the annotated API element  
 * is preliminary and subject to change.  
 */  
public @interface Preliminary { }
```

- Usage – No need to have ()
[@Preliminary](#)
public class TimeTravel { ... }

Single Value Annotation

- An annotation type with a single element
 - > The element should be named “**value**”

- Definition

```
/**
```

```
 * Associates a copyright notice with the annotated API element.
```

```
*/
```

```
public @interface Copyright {  
    String value();  
}
```

- Usage – can omit the element name and equals sign (=)
@Copyright("2002 Yoyodyne Propulsion Systems")
public class SomeClass { ... }

Normal Annotation

- We already have seen an example
- Definition

```
public @interface RequestForEnhancement {  
    int id();  
    String synopsis();  
    String engineer() default "[unassigned]";  
    String date(); default "[unimplemented]";  
}
```

- Usage

```
@RequestForEnhancement(  
    id = 2868724,  
    synopsis = "Enable time-travel",  
    engineer = "Mr. Peabody",  
    date = "4/1/3007"  
)  
public static void travelThroughTime(Date destination) { ... }
```

Annotation: **Meta-Annotations**

@Retention Meta-Annotation

- How long annotation information is kept
- Enum RetentionPolicy
 - > SOURCE - SOURCE indicates information will be placed in the source file but will not be available from the class files
 - > CLASS (Default)- CLASS indicates that information will be placed in the class file, but will not be available at runtime through reflection
 - > RUNTIME - RUNTIME indicates that information will be stored in the class file and made available at runtime through reflective APIs

@Target Meta-Annotation

- Restrictions on use of this annotation
- Enum ElementType
 - > TYPE, FIELD, METHOD, PARAMETER, CONSTRUCTOR, LOCAL_VARIABLE, ANNOTATION_TYPE, PACKAGE

Example: Definition and Usage of an Annotation with Meta Annotation

Definition of Accessor annotation

```
@Target(ElementType.FIELD)
@Retention(RetentionPolicy.CLASS)
public @interface Accessor {
    String variableName();
    String variableType() default "String";
}
```

Usage Example of the Accessor annotation

```
@Accessor(variableName = "name")
public String myVariable;
```

Reflection

- Check if MyClass is annotated with @Name annotation

```
boolean isName =  
    MyClass.class.isAnnotationPresent(Name.class);
```

Reflection

- Get annotation value of the @Copyright annotation

```
String copyright = MyClass.class.getAnnotation  
    (Copyright.class).value();
```

- Get annotation values of @Author annotation

```
Name author =  
    MyClass.class.getAnnotation(Author.class).value()  
String first = author.first();  
String last = author.last();
```

Concurrency

Concurrency Utilities: JSR-166

- Enables development of simple yet powerful multi-threaded applications
 - > Like Collection provides rich data structure handling capability
- Beat C performance in high-end server applications
 - > Fine-grained locking, multi-read single write lock
- Provide richer set of concurrency building blocks
 - > wait(), notify() and synchronized are too primitive
- Enhance scalability, performance, readability and thread safety of Java applications

Why Use Concurrency Utilities?

- Reduced programming effort
- Increased performance
- Increased reliability
 - > Eliminate threading hazards such as deadlock, starvation, race conditions, or excessive context switching are eliminated
- Improved maintainability
- Increased productivity

Concurrency Utilities

- Task Scheduling Framework
- Callable's and Future's
- Synchronizers
- Concurrent Collections
- Atomic Variables
- Locks
- Nanosecond-granularity timing

Concurrency: Task Scheduling Framework

Task Scheduling Framework

- **Executor/ExecutorsService/Executors** framework supports
 - > standardizing invocation
 - > scheduling
 - > execution
 - > control of asynchronous tasks according to a set of execution policies
- **Executor** is an interface
- **ExecutorsService** extends **Executor**
- **Executors** is factory class for creating various kinds of **ExecutorsService** implementations

Executor Interface

- **Executor** interface provides a way of de-coupling task **submission** from the **execution**
 - > execution: mechanics of how each task will be run, including details of thread use, scheduling

- Example

```
Executor executor = getSomeKindofExecutor();  
executor.execute(new RunnableTask1());  
executor.execute(new RunnableTask2());
```

- Many **Executor** implementations impose some sort of limitation on how and when tasks are scheduled

Executor and ExecutorService

ExecutorService adds lifecycle management

```
public interface Executor {  
    void execute(Runnable command);  
}  
  
public interface ExecutorService extends Executor {  
    void shutdown();  
    List<Runnable> shutdownNow();  
    boolean isShutdown();  
    boolean isTerminated();  
    boolean awaitTermination(long timeout,  
                             TimeUnit unit);  
  
    // other convenience methods for submitting tasks  
}
```

Creating ExecutorService From Executors

```
public class Executors {  
    static ExecutorService  
        newSingleThreadedExecutor();  
  
    static ExecutorService  
        newFixedThreadPool(int n);  
  
    static ExecutorService  
        newCachedThreadPool(int n);  
  
    static ScheduledExecutorService  
        newScheduledThreadPool(int n);  
  
    // additional versions specifying ThreadFactory  
    // additional utility methods  
}
```

pre-J2SE 5.0 Code

Web Server—poor resource management

```
class WebServer {  
  
    public static void main(String[] args) {  
        ServerSocket socket = new ServerSocket(80);  
  
        while (true) {  
            final Socket connection = socket.accept();  
            Runnable r = new Runnable() {  
                public void run() {  
                    handleRequest(connection);  
                }  
            };  
            // Don't do this!  
            new Thread(r).start();  
        }  
    }  
}
```

Executors Example

Web Server—better resource management

```
class WebServer {
    Executor pool =
        Executors.newFixedThreadPool(7);

    public static void main(String[] args) {
        ServerSocket socket = new ServerSocket(80);

        while (true) {
            final Socket connection = socket.accept();
            Runnable r = new Runnable() {
                public void run() {
                    handleRequest(connection);
                }
            };
            pool.execute(r);
        }
    }
}
```

Concurrency: Callables and Futures

Callable's and Future's: Problem (pre-J2SE 5.0)

- If a new thread (callable thread) is started in an application, there is currently no way to return a result from that thread to the thread (calling thread) that started it without the use of a shared variable and appropriate synchronization
 - > This is complex and makes code harder to understand and maintain

Callables and Futures

- Callable thread (Callee) implements **Callable** interface
 - > Implement **call()** method rather than **run()**
- Calling thread (Caller) submits **Callable** object to Executor and then moves on
 - > Through **submit()** not **execute()**
 - > The **submit()** returns a **Future** object
- Calling thread (Caller) then retrieves the result using **get()** method of **Future** object
 - > If result is ready, it is returned
 - > If result is not ready, calling thread will block

Build CallableExample (This is Callee)

```
class CallableExample
    implements Callable<String> {

    public String call() {
        String result = "The work is ended";

        /* Do some work and create a result */

        return result;
    }
}
```

Future Example (Caller)

```
ExecutorService es =  
    Executors.newSingleThreadExecutor();  
  
Future<String> f =  
    es.submit(new CallableExample());  
  
/*  Do some work in parallel  */  
  
try {  
    String callableResult = f.get();  
} catch (InterruptedException ie) {  
    /*  Handle  */  
} catch (ExecutionException ee) {  
    /*  Handle  */  
}
```

Concurrency:
Synchronizers: Semaphore

Semaphores

- Typically used to restrict access to fixed size pool of resources
- New Semaphore object is created with same count as number of resources
- Thread trying to access resource calls **acquire()**
 - > Returns immediately if semaphore count > 0
 - > Blocks if count is zero until **release()** is called by different thread
 - > **acquire()** and **release()** are thread safe atomic operations

Semaphore Example

```
private Semaphore available;  
private Resource[] resources;  
private boolean[] used;  
  
public Resource(int poolSize) {  
    available = new Semaphore(poolSize);  
    /* Initialise resource pool */  
}  
public Resource getResource() {  
    try { available.acquire() } catch (IE) {}  
    /* Acquire resource */  
}  
public void returnResource(Resource r) {  
    /* Return resource to pool */  
    available.release();  
}
```

Concurrency: Concurrent Collections

BlockingQueue Interface

- Provides thread safe way for multiple threads to manipulate collection
- **ArrayBlockingQueue** is simplest concrete implementation
- Full set of methods
 - > **put ()**
 - > **offer ()** [non-blocking]
 - > **peek ()** [non-blocking]
 - > **take ()**
 - > **poll ()** [non-blocking and fixed time blocking]

Blocking Queue Example: Logger placing log messages

```
private ArrayBlockingQueue messageQueue =  
    new ArrayBlockingQueue<String>(10);  
  
Logger logger = new Logger(messageQueue);  
  
public void run() {  
    String someMessage;  
    try {  
        while (true) {  
            /* Do some processing */  
  
            /* Blocks if no space available */  
            messageQueue.put(someMessage);  
        }  
    } catch (InterruptedException ie) { }  
}
```

Blocking Queue Example: Log Reader reading log messages

```
private BlockingQueue<String> msgQueue;

public LogReader(BlockingQueue<String> mq) {
    msgQueue = mq;
}

public void run() {
    try {
        while (true) {
            String message = msgQueue.take();
            /* Log message */
        }
    } catch (InterruptedException ie) {
        /* Handle */
    }
}
```

Concurrency: **Atomic Variables**

Atomics

- `java.util.concurrent.atomic`
 - > Small toolkit of classes that support lock-free thread-safe programming on single variables

```
AtomicInteger balance = new AtomicInteger(0);
```

```
public int deposit(integer amount) {  
    return balance.addAndGet(amount);  
}
```

Concurrency: Locks

Locks

- Lock interface
 - > More extensive locking operations than synchronized block
 - > Caution: No automatic unlocking like synchronized block – use try/finally to unlock
 - > Advantage: Non-blocking access is possible using `tryLock()`
- ReentrantLock
 - > Concrete implementation of Lock
 - > Holding thread can call `lock()` multiple times and not block
 - > Useful for recursive code

ReadWriteLock

- Has two locks controlling read and write access
 - > Multiple threads can acquire the read lock if no threads have a write lock
 - > If a thread has a read lock, others can acquire read lock but nobody can acquire write lock
 - > If a thread has a write lock, nobody can have read/write lock
 - > Methods to access locks

```
rw1.readLock().lock();  
rw1.writeLock().lock();
```

ReadWrite Lock Example

```
class ReadWriteMap {
    final Map<String, Data> m = new TreeMap<String, Data>();
    final ReentrantReadWriteLock rwl =
        new ReentrantReadWriteLock();
    final Lock r = rwl.readLock();
    final Lock w = rwl.writeLock();
    public Data get(String key) {
        r.lock();
        try { return m.get(key) }
        finally { r.unlock(); }
    }
    public Data put(String key, Data value) {
        w.lock();
        try { return m.put(key, value); }
        finally { w.unlock(); }
    }
    public void clear() {
        w.lock();
        try { m.clear(); }
        finally { w.unlock(); }
    }
}
```


JMX (Java Management Extension)

JMX Introduction

- Overview of JMX
- Instrument your Application
- Accessing your instrumentation remotely
- What's coming in JDK 6

What is JMX?

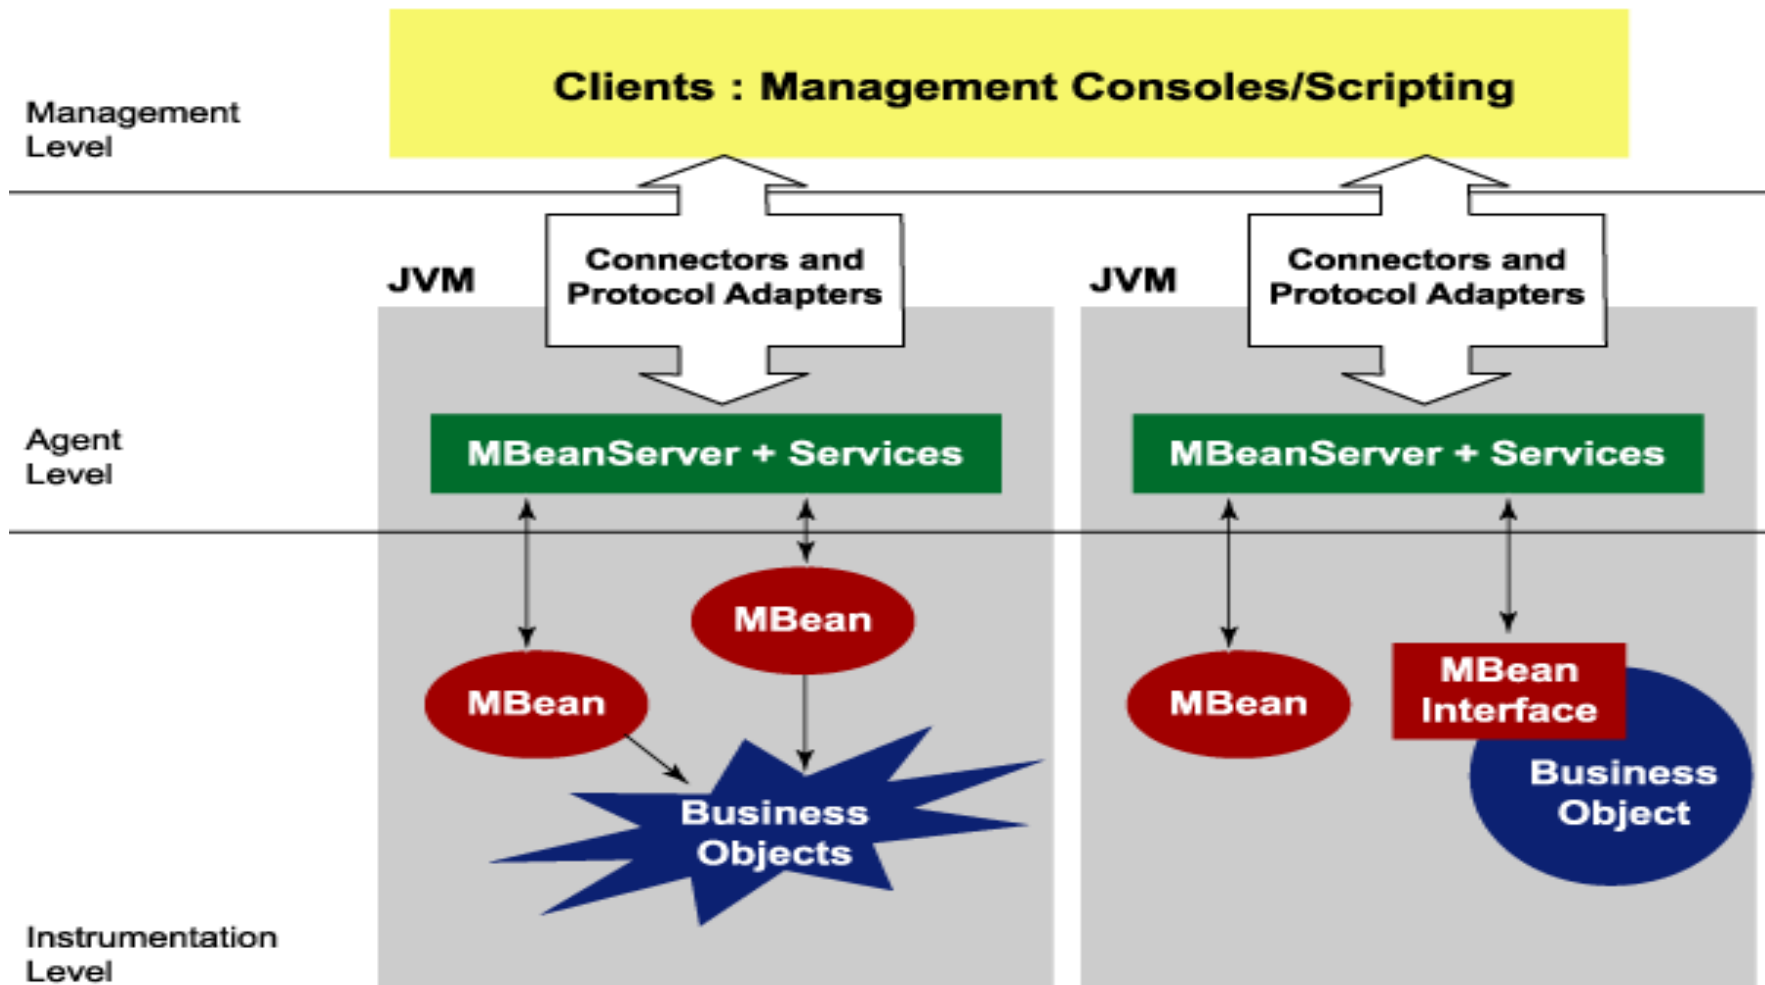
- Standard API for developing observable applications – JSR 3 and JSR 160
- Provides access to information such as
 - > Number of classes loaded
 - > Virtual machine uptime
 - > Operating system information
- Applications can use JMX for
 - > Management – changing configuration settings
 - > Monitoring – getting statistics and notifications
- Mandatory in J2SE 5.0 and J2EE 1.4

JMX: Architecture

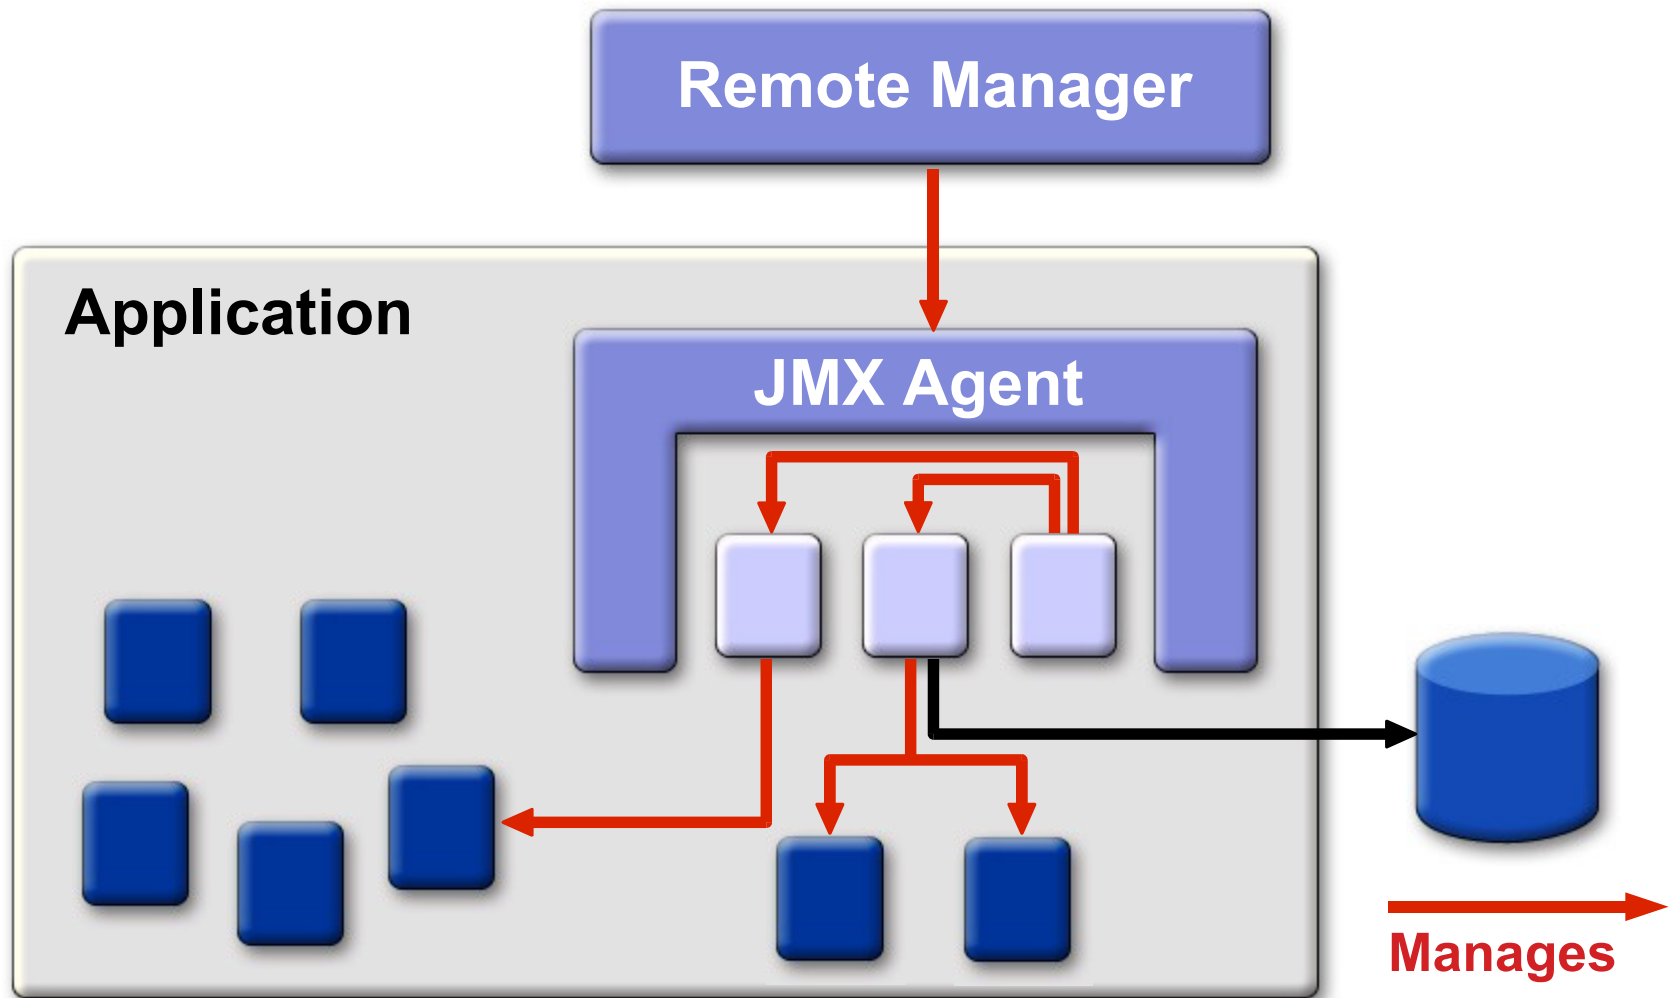
JMX Architecture

- Instrumentation Level
 - > MBeans instruments resources, exposing attributes and operations
- Agent Level
 - > MBean Server
 - > Predefined services
- Remote Management
 - > Protocol Adaptors and Standard Connectors enables remote Manager Applications

JMX Architecture



JMX Architecture



JMX: MBean

Managed Beans(MBeans)

- A MBean is a named *managed object* representing a *resource*
 - > An application configuration setting
 - > Device
 - > Etc.
- A MBean can have
 - > Attributes that can be read and/or written
 - > Operations that can be invoked
 - > Notifications that the MBean can broadcast

A MBean Example

CacheControlMBean			
Used: int	R	}	attributes
Size: int	RW		
save(): void		}	operations
dropOldest(int n): int			
"com.example.config.change"		}	notifications
"com.example.cache.full"			

Standard MBean

- Standard MBean is the simplest model to use
 - > Quickest and Easiest way to instrument static manageable resources
- Steps to create a standard MBean
 - > Create an Java interface call **FredMBean**
 - > Follows JavaBeans naming convention
 - > Implement the interface in a class call **Fred**
- An instance of **Fred** is the MBean

Dynamic MBean

- Expose attributes and operations at **Runtime**
- Provides more flexible instrumentations
- Step to create Dynamic MBeans
 - > Implements **DynamicMBeans** interface
 - > Method returns all Attributes & Operations
- The same capability as Standard MBeans from Agent's perspective

DynamicMBean Interface

**<<Interface>>
DynamicMBean**

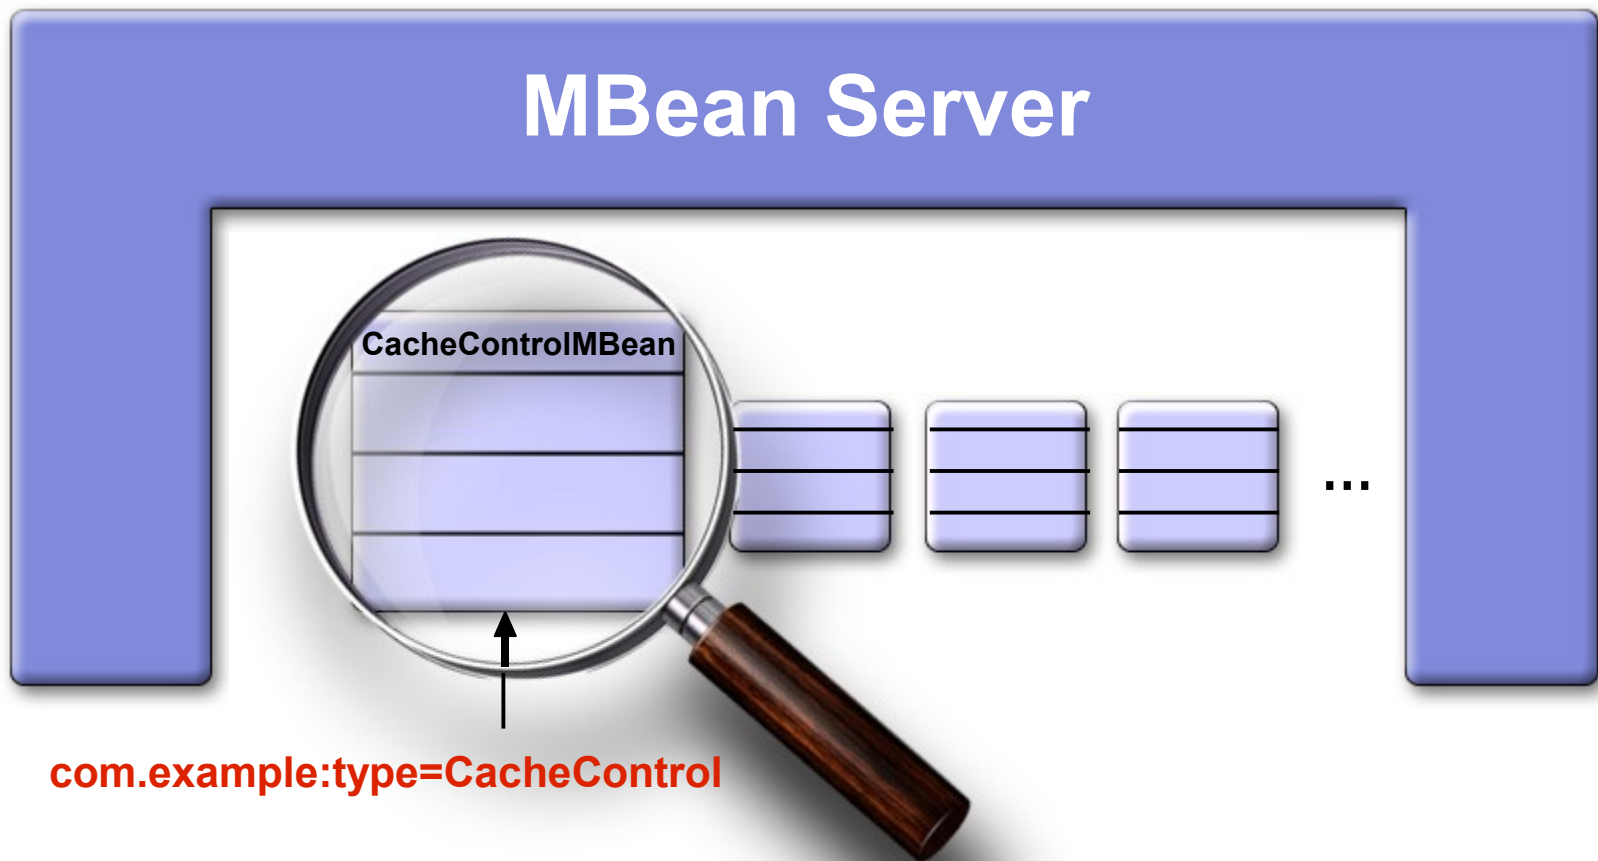
**getMBeanInfo():MBeanInfo
getAttribute(attribute:String):Object
getAttributes(attributes:String[]):AttributeList
setAttribute(attribute:Attribute):void
setAttributes(attributes:AttributeList):AttributeList
invoke(actionName:String,
 params:Object[],
 signature:String[]):Object**

JMX Notification

- JMX notifications consists of the following
 - > **NotificationEmitter** – event generator, typically your MBean
 - > **NotificationListener** – event listener
 - > **Notification** – the event
 - > **NotificationBroadcasterSupport** – helper class
- Register with MBean server to receive events

JMX: MBean Server

MBean Server




MBean Server

- To be useful, an MBean must be registered in an MBean Server
- Usually, the only access to MBeans is through the MBean Server
- You can have more than one MBean Server per Java™ Virtual Machine (JVM™ machine)
- But usually, as of Java SE 5, everyone uses the Platform MBean Server
 - > `java.lang.management.ManagementFactory.
getPlatformMBeanServer()`

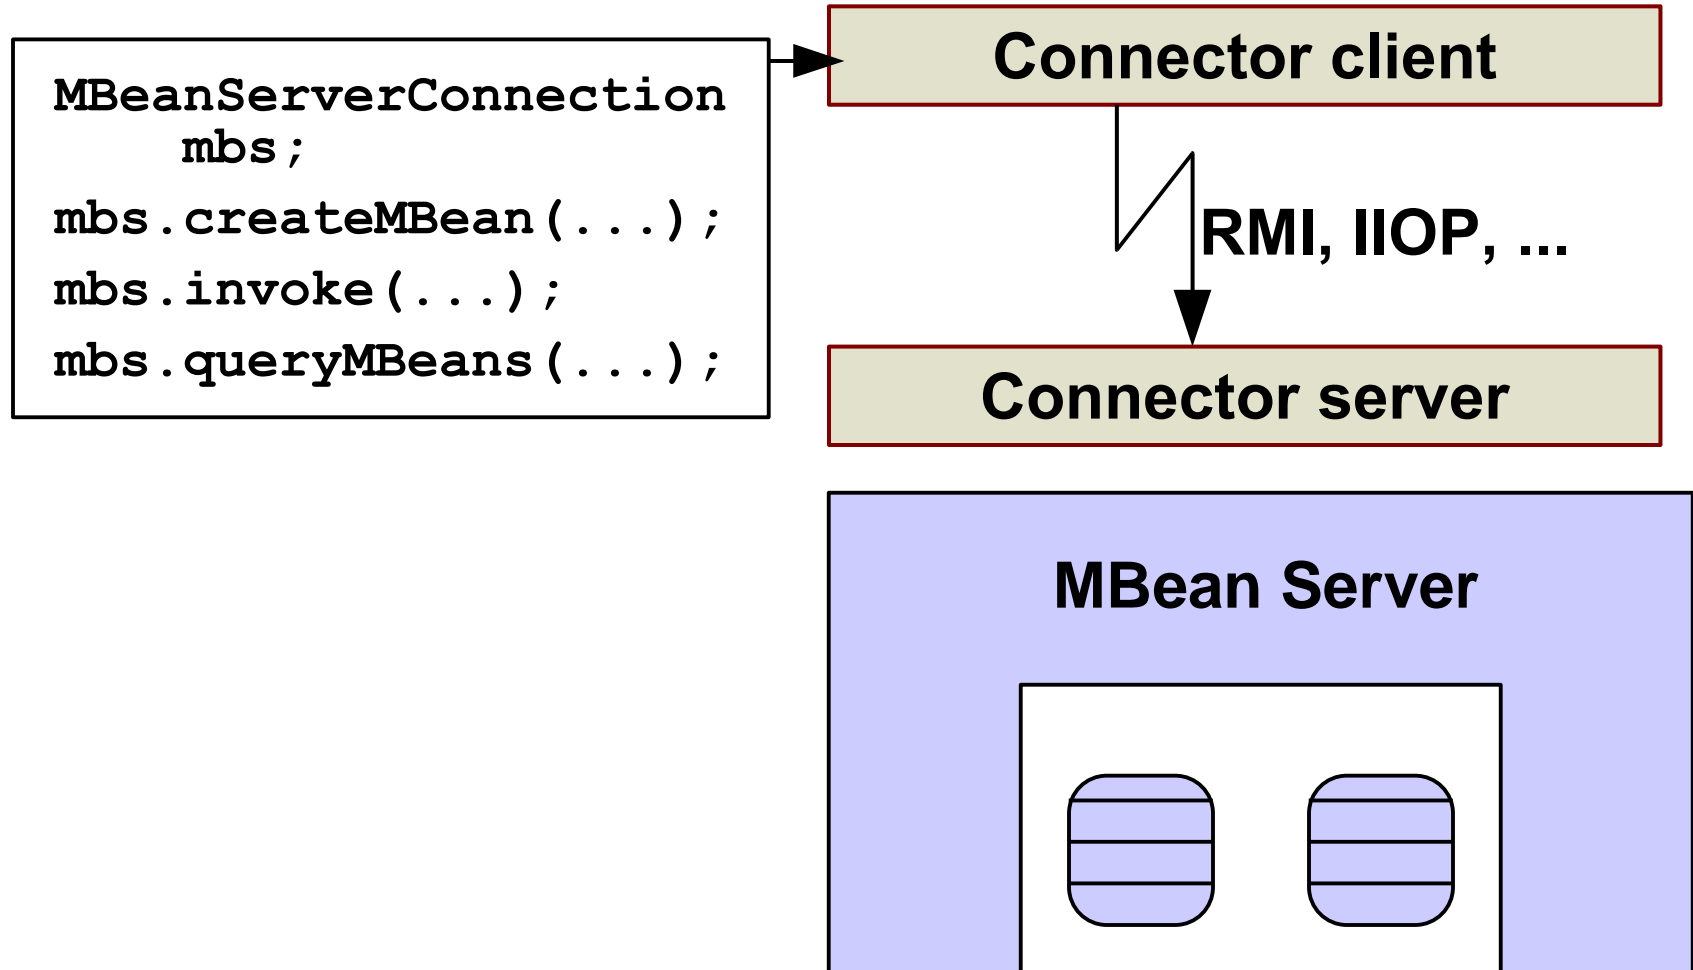
JMX: Client Types

MBean Server: Local Clients

```
MBeanServer mbs;  
  
mbs.createMBean(...);  
mbs.invoke(...);  
mbs.queryMBeans(...);
```



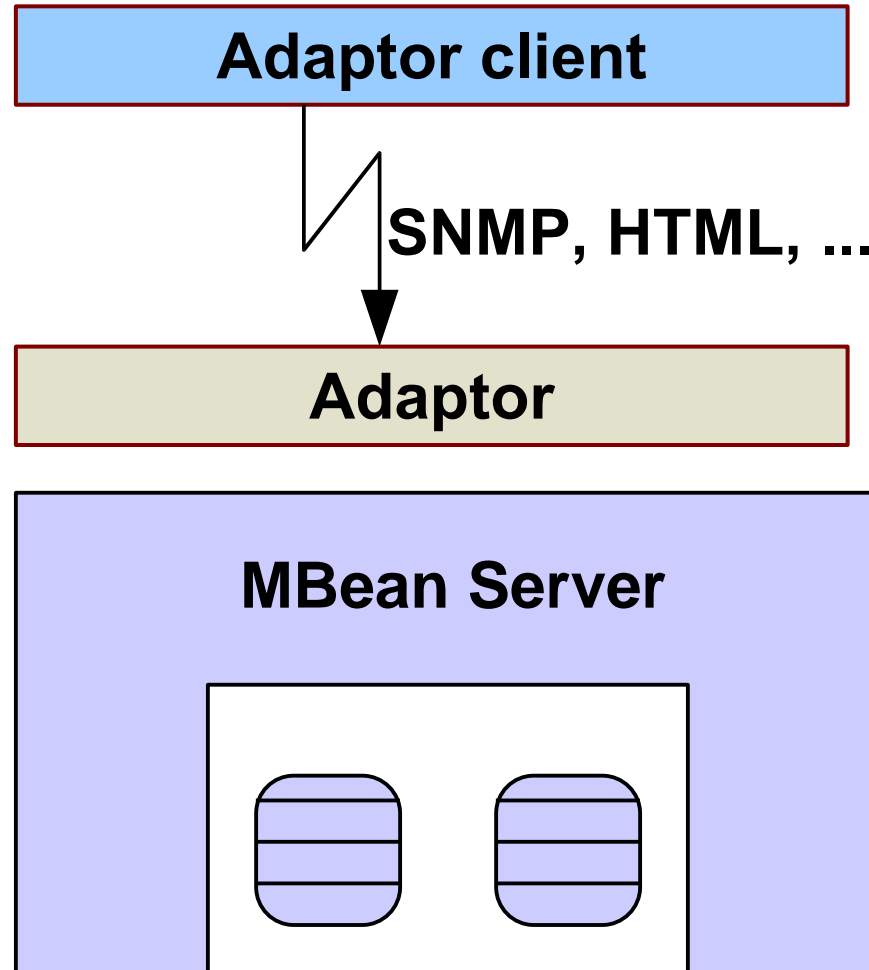
MBean Server: Connector Client




MBean Server: Connector

- Connectors defined by the JMX Remote API (JSR 160)
 - > Unrelated to the J2EE™ Connector Architecture
- Java SE architecture includes RMI and RMI/IIOP connectors
- JSR 160 also defines a purpose-built protocol, JMXMP
- Future work: a SOAP-based connector for the Web Services world (JSR 262)

MBean Server: Adaptor Client



Mapping SNMP or CIM to JMX API



- Generation not currently standard
 - > proprietary solutions exist (Sun's is JDMK)
- Implementing semantics may mean mapping to another, “native” JMX API model
- Automated reverse mapping from JMX API to SNMP or CIM gives poor results

JMX: JMX API Services

JMX API Services

- JMX API includes a number of pre-defined services
 - > Services are themselves MBeans
- Monitoring service (thresholding)
 - > javax.management.monitor
- Relation service (relations between MBeans)
 - > javax.management.relation
- Timer service
 - > javax.management.timer
- M-let service
 - > javax.management.loading

JMX: **Steps of instrumenting Your Application**

Steps for Instrumenting Your App

- Create MBean's
 - > Define an MBean interface
 - > Add attributes and operations
 - > Add notifications
 - > Implement MBean interface
- Create JMX agent
 - > Provides a method to create and register your MBeans.
 - > Provides access to the MBean server
- Run the application with JConsole

JMX:

Demo – Running Anagram application with JMX support

Demo Scenario

- Anagram game is managed via JMX
 - > Manage and monitor number of seconds it takes a user to provide a right answer
 - > Monitor number of times a user has provided solutions
 - > Subscribe event notification



Java SE Beyond Basics: Generics, Annotation, JMX

Sang Shin

sang.shin@sun.com

www.javapassion.com

Sun Microsystems Inc.

**UNLOCK
OPPORTUNITY**

What will you open?



SUN TECH DAYS 2006-2007
A Worldwide Developer Conference

Instrument ClickCounter (Standard MBean)

- Create ClickCounterStdMBean interface
- Create Mbean ClickCounterStd implementing ClickCounterStdMBean, and extending ClickFrame as well
- Get System MBean Server
- Register ClickCounterStd
- We're done!

Accessing the JMX Agent

- J2SE 5.0 and later releases
 - > `java -Dcom.sun.management.jmxremote Main`
 - > See <http://java.sun.com/j2se/1.5.0/docs/guide/management/agent.html> for more options
- Start `jconsole`
- In JavaSE 6.0, you can attach `jconsole` without setting `com.sun.management.jmxremote` property
- One click monitoring with NetBeans IDE

Standard Mbean Interface

```
public interface ClickCounterStdMBean {  
    public void reset();  
    public int getDisplayNumber();  
    public void setDisplayNumber(int inNumber);  
    public int getCountNumber();  
}
```

Implements Standard MBean

```
public class ClickCounterStd
    extends ClickFrame
    implements ClickCounterStdMBean {

    public void reset() {
        getModel().reset();
        updateLabel();
    }

    public int getDisplayNumber() {
        return getModel().getDisplayNumber();
    }
    . . . . .
}
```

Registering a MBean

```
MBeanServer mbs = ManagementFactory  
    .getPlatformMBeanServer( );
```

```
ObjectName name = new ObjectName(  
    "shen.joey.demo.ClickCount:type=ClickCounterStd" );
```

```
ClickFrame counter = new ClickCounterStd( );
```

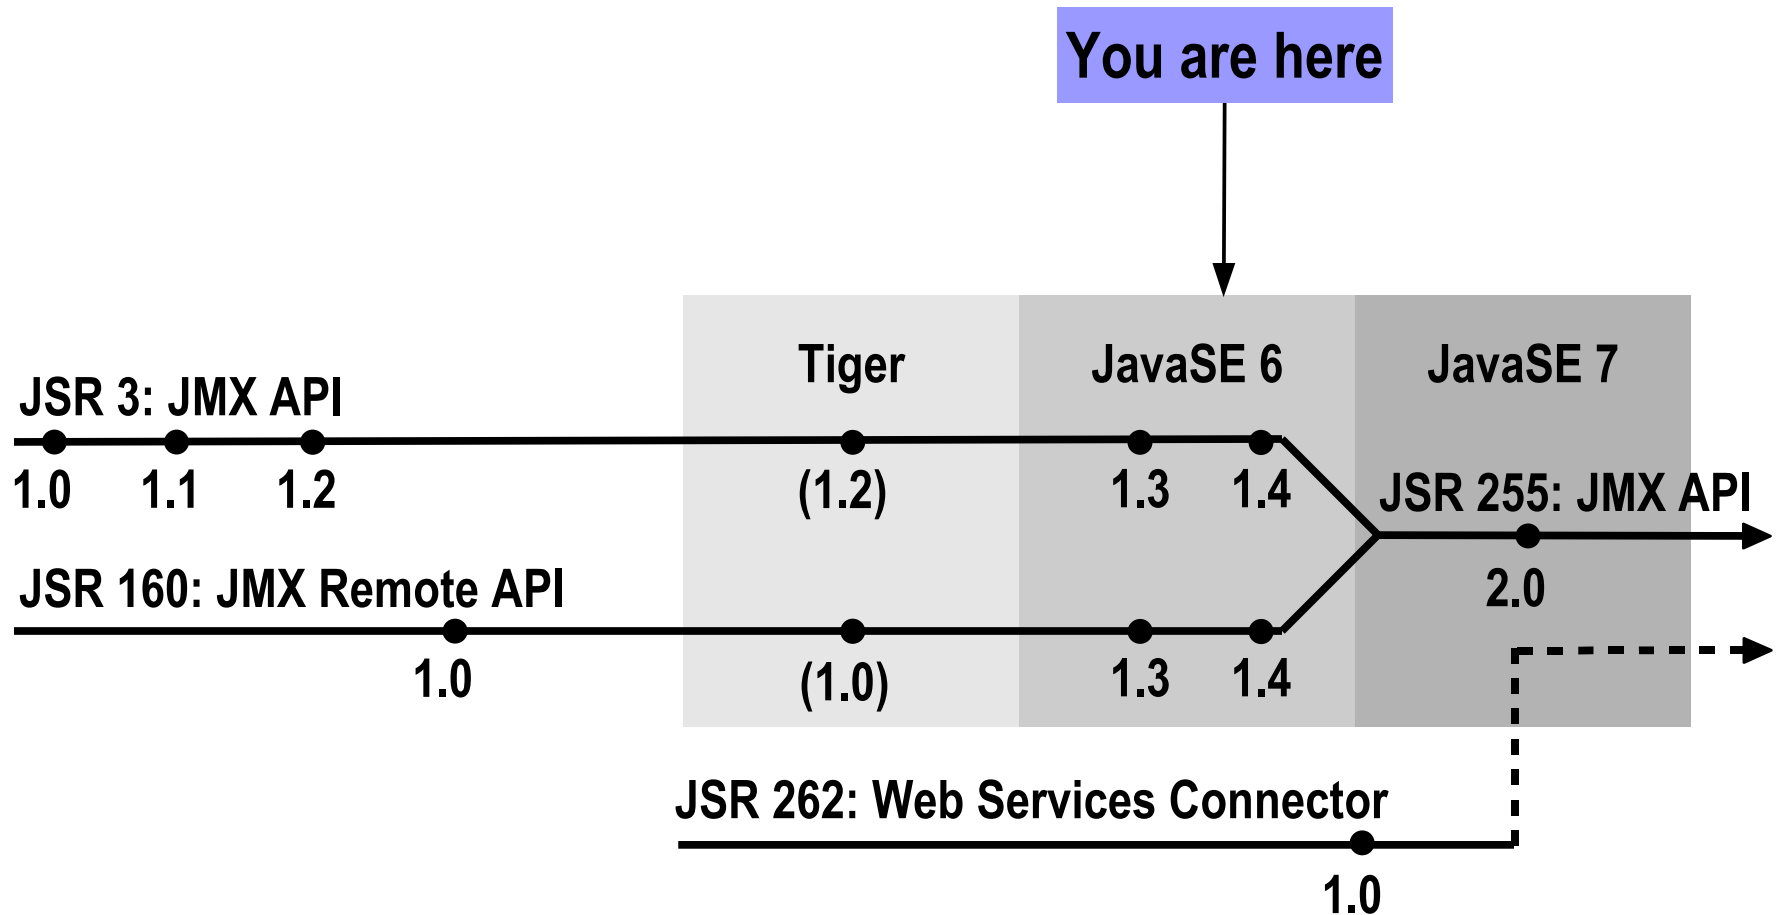
```
mbs.registerMBean(counter, name);
```

JMX: Roadmap

What's coming next?

- JSR 3 defined the JMX API
 - > Updated in Java Platform, Standard Edition 6
- JSR 160 defined the JMX Remote API
 - > Also updated in JavaSE 6.0
- JSR 255 merges and updates JSRs 3 and 160
 - > It will produce JMX API version 2.0 in Java SE 7

JMX API Versions



<http://jdk6.dev.java.net/>

MXBeans: Problem Statement (1)

- An MBean interface can include arbitrary Java programming language types

```
public interface ThreadMBean {  
    public ThreadInfo getThreadInfo();  
}  
  
public class ThreadInfo {  
    public String getName();  
    public long getBlockedCount();  
    public boolean isSuspended();  
    ...  
}
```

- When values must be grouped automatically

MXBeans: Problem Statement (2)

- An MBean interface can include arbitrary Java programming language types

```
public interface ThreadMBean {  
    public ThreadInfo getThreadInfo();  
}
```

- Client must have these classes
- What about generic clients like jconsole?
- What about versioning?

MXBeans (1)

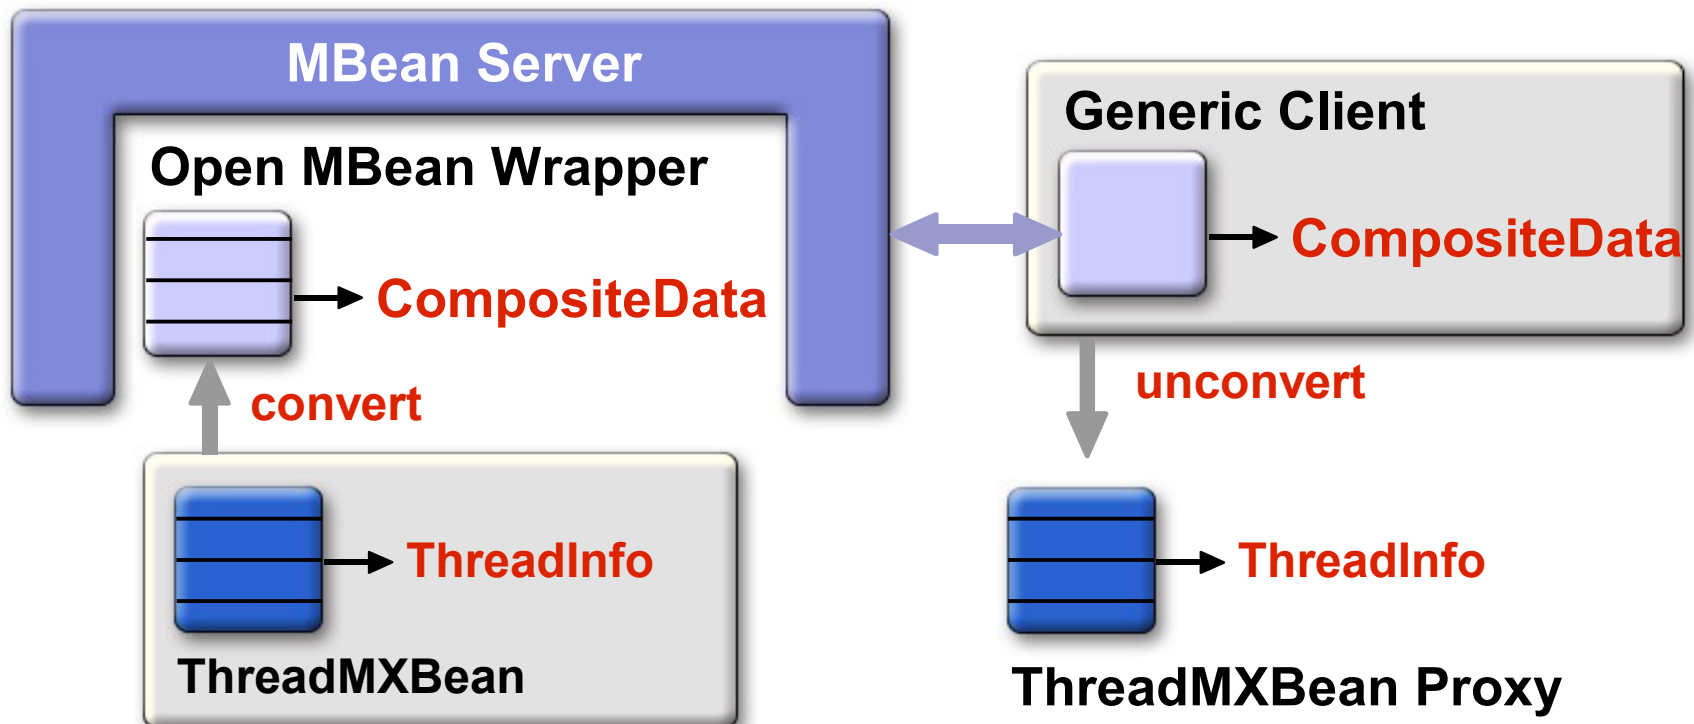
- MXBeans were designed for the instrumentation of the VM itself (JSR 174)
 - > Already exist in `java.lang.management`
 - > User-defined MXBeans are new in Mustang
- Management interface still a bean interface
- Can reference arbitrary types, with some restrictions

MXBeans (2)

```
public interface ThreadMXBean {  
    public ThreadInfo getThreadInfo();  
}  
public class ThreadMXBeanImpl implements ThreadMXBean {  
    // Do not need Something/SomethingMXBean naming  
    public ThreadInfo getThreadInfo() {  
        return new ThreadInfo(...);  
    }  
}  
ThreadMXBean mxbean = new ThreadMXBeanImpl();  
ObjectName name =  
    new ObjectName("java.lang:type=Threading");  
  
mbs.registerMBean(mxbean, name);
```

MXBeans (3)


- Generic client can access as Open MBean
- Model-aware client can make ThreadMXBean proxy



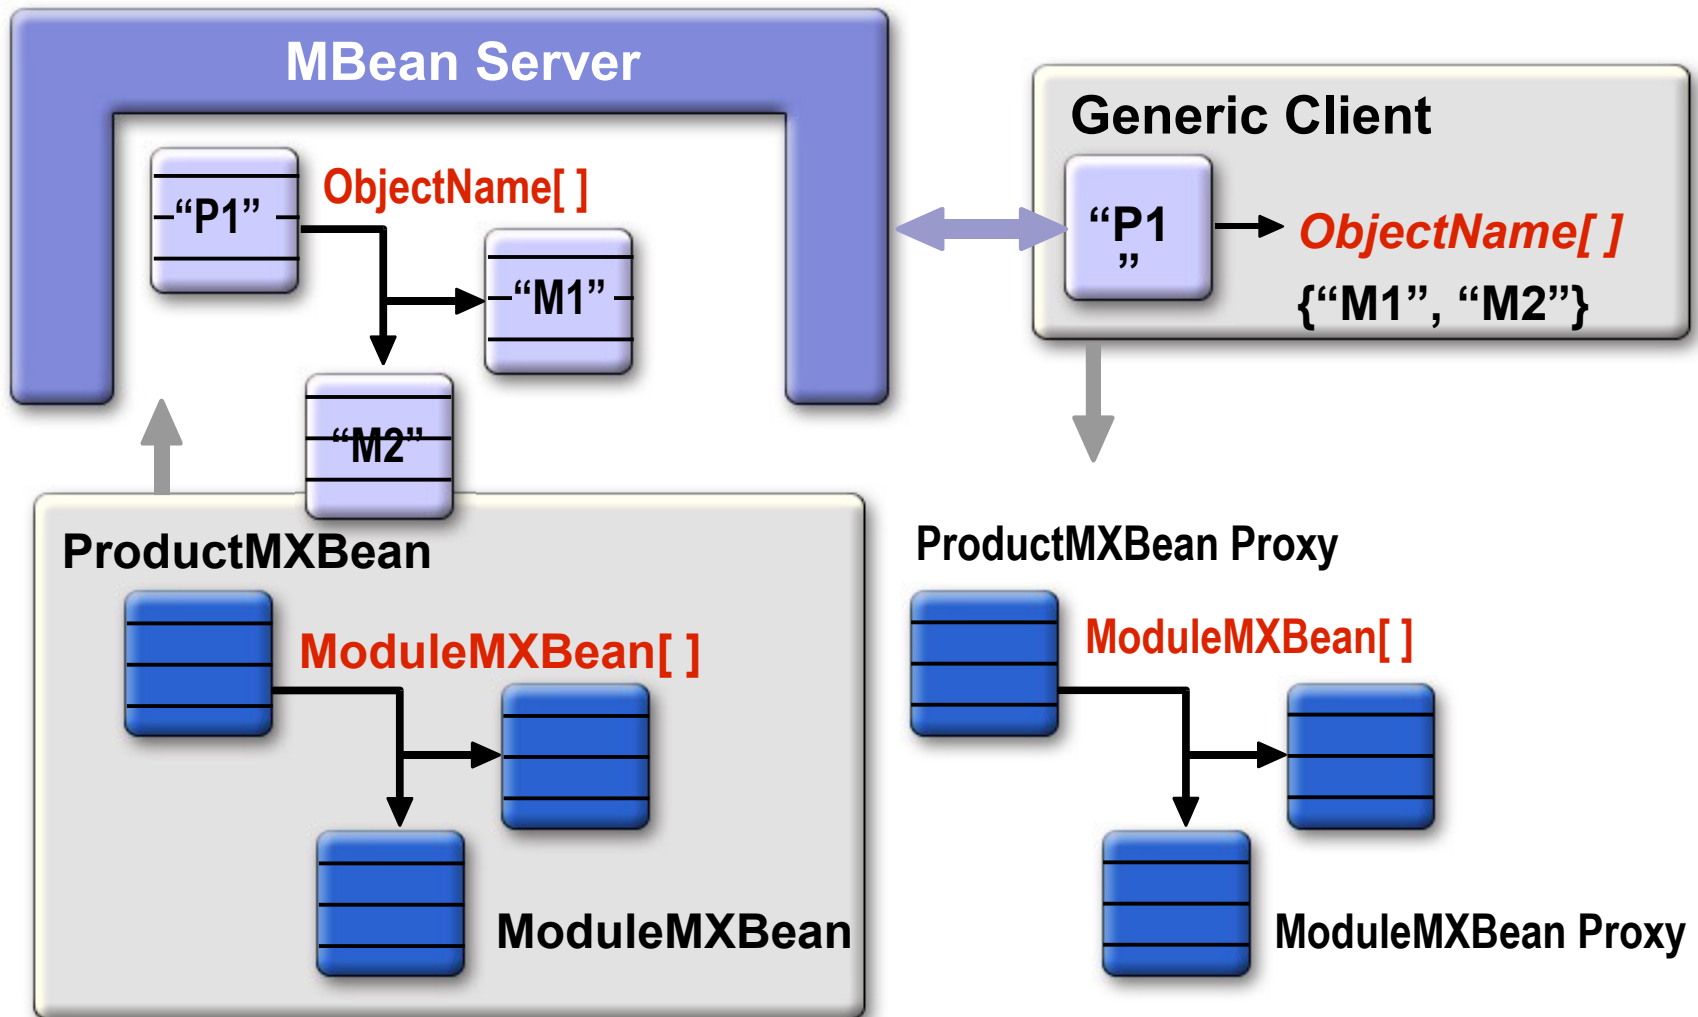
MXBean References (1)

```
public interface ProductMXBean {
    ModuleMXBean[] getModules();
}
```

ProductMXBean



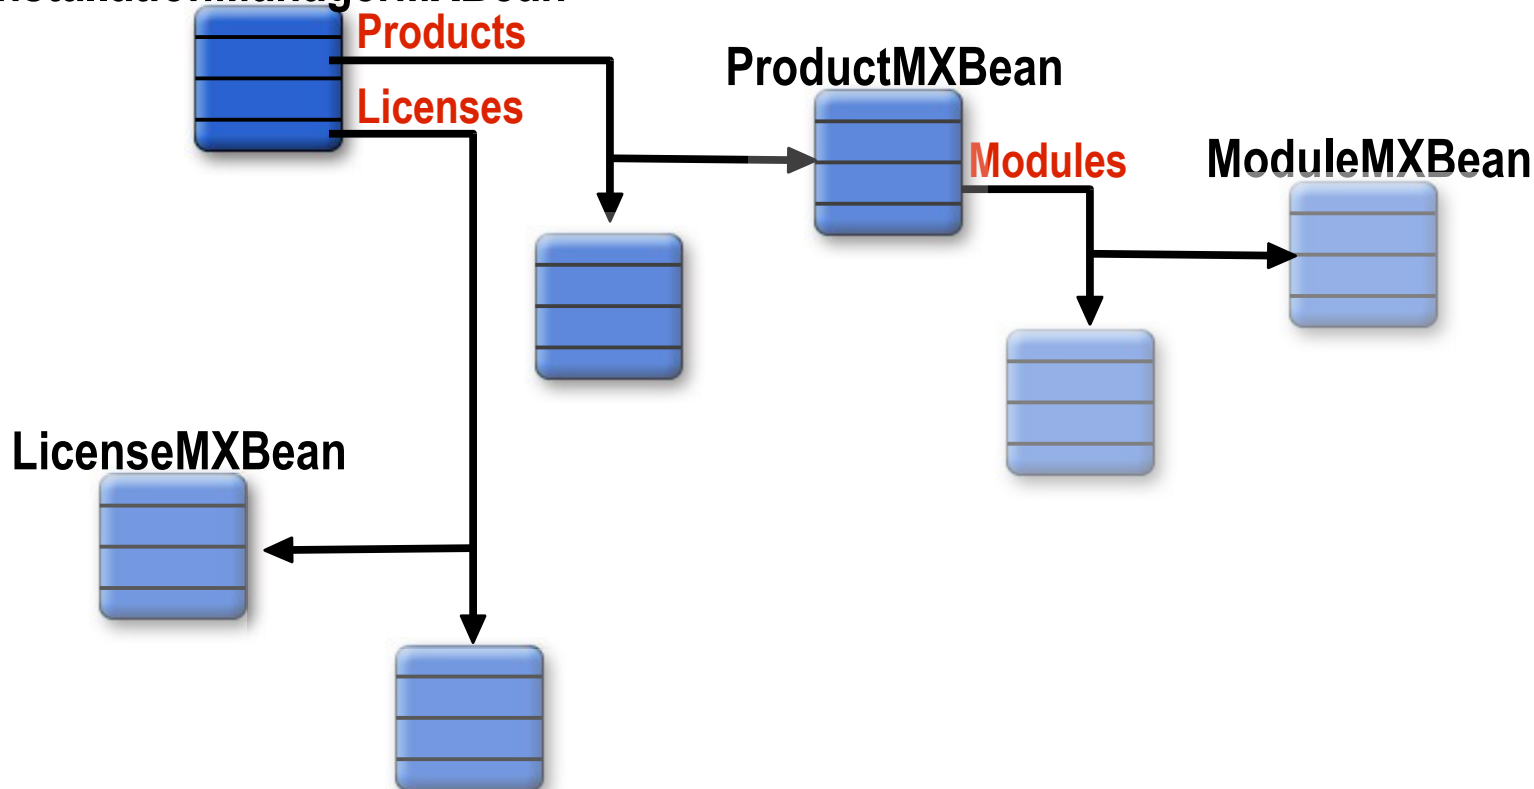
MXBean References (2)



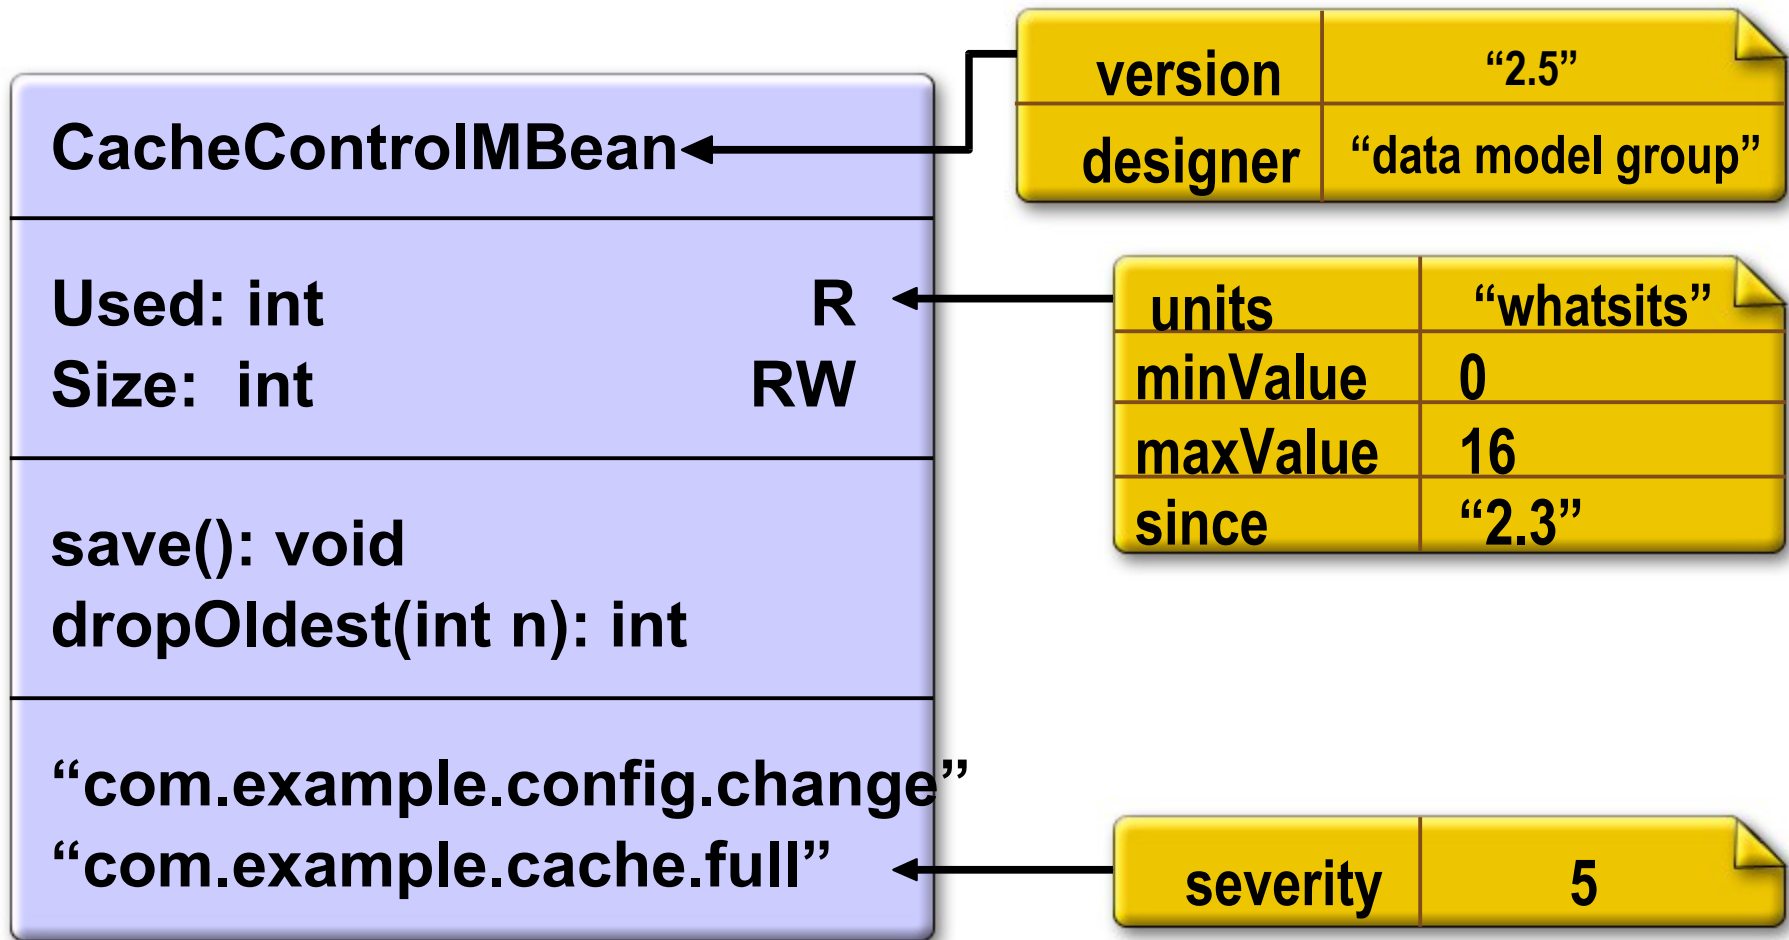
MXBean References (3)

Navigating From a Starting Point

InstallationManagerMXBean



Descriptors




Descriptors and Generic Clients

(like jconsole)

Used: int R

units	"whatsits"
minValue	0
maxValue	16

Used



Descriptor Details

- Classes MBeanInfo, MBeanAttributeInfo, etc., now have an optional Descriptor
- Every attribute, operation, notification can have its own Descriptor
- Descriptor is set of (key,value) pairs
- Some keys have conventional meanings
- Users can add their own keys
- Descriptors have always existed in Model MBeans

Descriptor Annotations

```
public interface CacheControlMBean {
    @Units("whatsits") @Range(minValue=0, maxValue=16)
    public int getUsed();
}
```

Used: int

R

units	"whatsits"
minValue	0
maxValue	16

With definitions like:

```
public @interface Range {
    @DescriptorKey("minValue")
    public int minValue();
    @DescriptorKey("maxValue")
    public int maxValue();
}
```

Some Other Mustang changes

- Generified at last!
 - > Set<ObjectName> queryNames(...)
- More-general ObjectName wildcards
 - > domain:type=Dir,path="/root/ * "
- Simpler Notification use
 - > NotificationBroadcasterSupport(MBeanNotificationInfo[])
 - > class StandardEmitterMBean extends StandardMBean
- Monitor attributes of complex type
 - > MonitorMBean.setObservedAttribute("ThreadInfo.size")

<http://jdk6.dev.java.net>

JMX Summary

- JMX core concept
 - > MBean, MXBean, MBeanServer
- Instrument App with JMX
 - > Create MBean
 - > Register to MBeanServer
- Coming Next
 - > MXBean
 - > Descriptor

Performance

GC Algorithms

- Reference Counting
 - > It's straightforward and easy
 - > Need support of compilers
 - > Circularly referenced detection
- Tracing collectors
 - > Root Objects: Local Variable, Static Variable, Registers

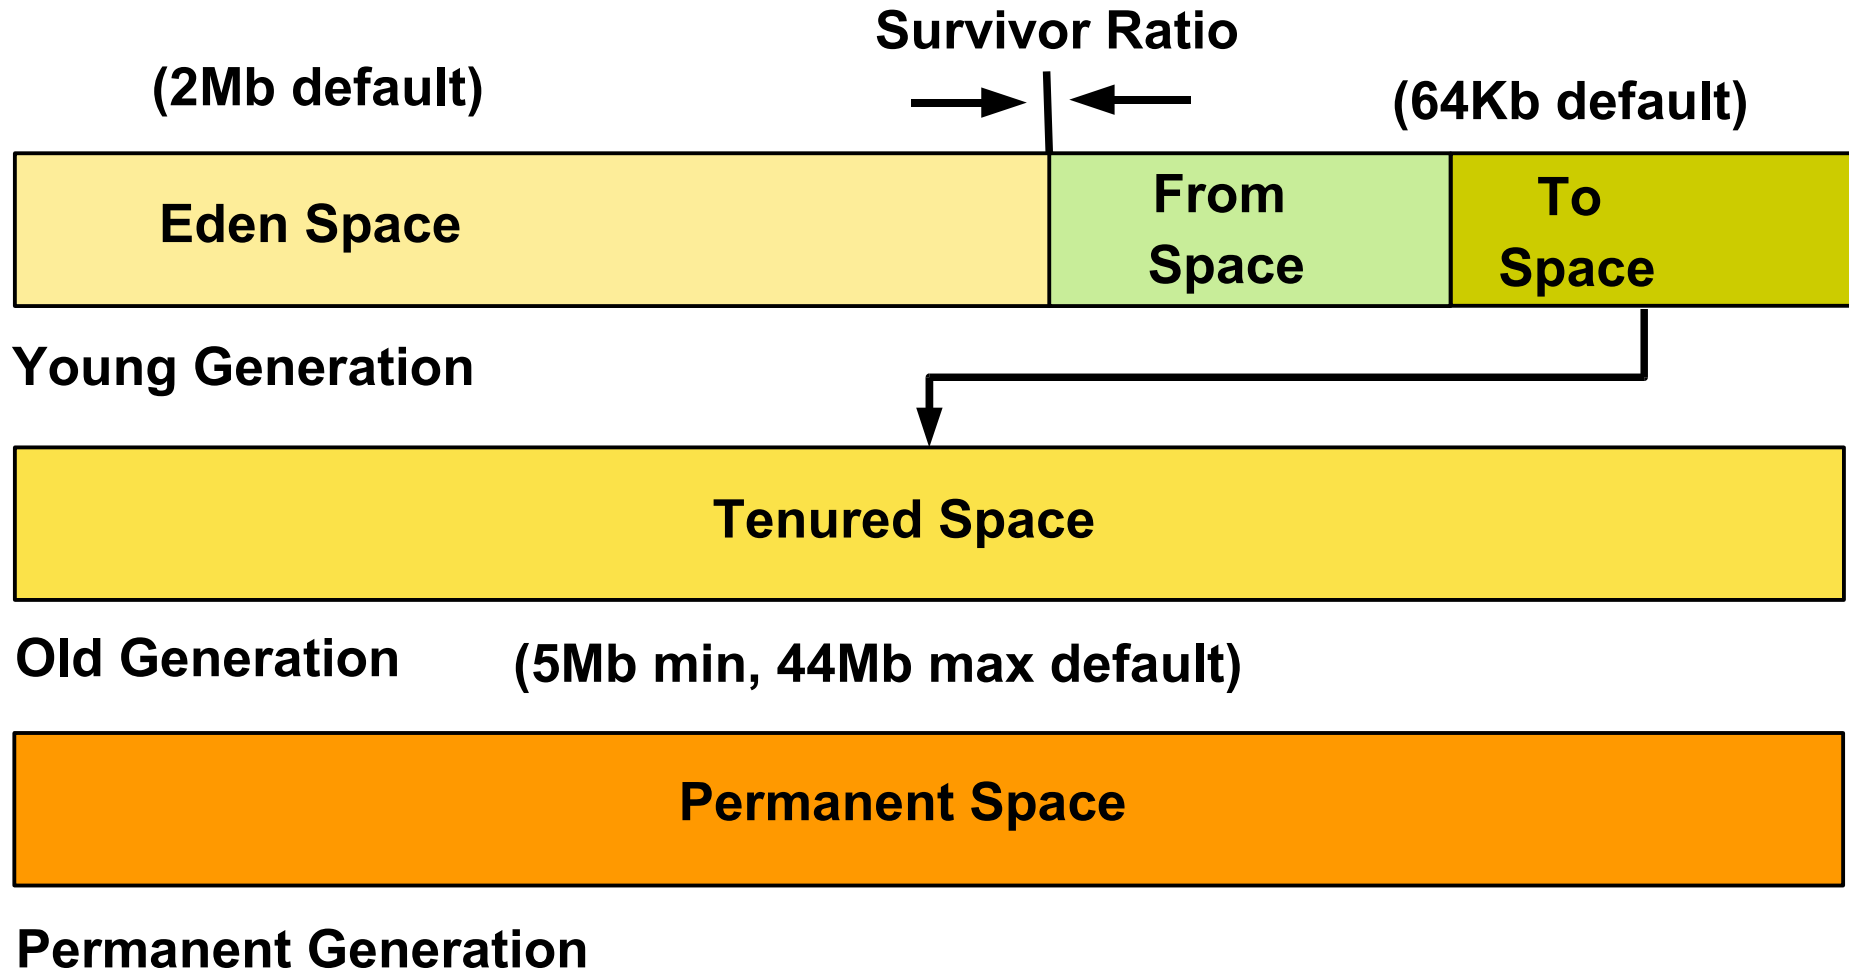
Mark-Sweep Collector

- Behavior
 - > Stop the world
 - > Mark all the reachable objects
 - > Exam all the heap, Sweep the unreachable objects
- Advantage
 - > Simple to implement
 - > Dose not depend on compilers
- Limitation
 - > Every allocated object is visited
 - > Heap fragment

Copying Collector

- Advantage
 - > Visit only live objects
 - > Compact in the new space
 - > Greatly reduces the cost of object allocation
 - > Easy reclaim
- Limitation
 - > Larger memory footprint
 - > Overhead of copying
 - > Long live objects
 - > Adjust reference address

HotSpot Heap Layout



General Tuning Advice

- Allocate as much memory as possible to VM (As long as pause time is not the problem)
 - > 64M default is often too small
- Set -Xms and -Xmx the same
 - > Increase predictability, improve startup time
- Set Eden/Tenured ratio
 - > Eden < 50% (Not for throughput and Concurrent collectors)
 - > NewRatio=2 seems to be good
- Disable explicit GC
 - > -XX:+DisableExplicitGC

Throughput Collector

- When there are a large number of processors
- Parallel version of the young generation collector
- -XX:+UseParallelGC to enable
- J2SE 1.5 will automatically choose throughput collector
- on Server Machines(2+ Processors, 2G+ Memory)
- -XX:MaxGCPauseMillis=<nnn>
- -XX:GCTimeRatio=<Apps time/GC time>

Concurrent Collector

- For the sack of low pause time
- Applications which have a large set of long-lived data running on more than one processor
- Parallel in Young generation collecting, Concurrent in Tenured generation collecting
- `-XX:+UseConcMarkSweepGC` to enable

Resources

- <http://java.sun.com/j2se/1.5.0/docs/guide/jmx/tutorial/tutorialTOC.html>
- <http://java.sun.com/docs/performance/>
- <http://java.sun.com/j2se/1.5.0/docs/guide/concurrency/index.html>